

# Recurrent Neural Networks

## Remembering Past Data

Dr. Mauricio Toledo-Acosta

Diplomado Ciencia de Datos con Python

# Table of Contents

1 Recurrent Neural Networks

2 LSTM

# Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.

# Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form  $x_1, \dots, x_n$ , where  $x_t$  is a  $d$ -dimensional point received at the time-stamp  $t$ .

# Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form  $x_1, \dots, x_n$ , where  $x_t$  is a  $d$ -dimensional point received at the time-stamp  $t$ .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input  $x_t$  only after the earlier inputs have already been received and converted into a hidden state.

# Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form  $x_1, \dots, x_n$ , where  $x_t$  is a  $d$ -dimensional point received at the time-stamp  $t$ .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input  $x_t$  only after the earlier inputs have already been received and converted into a hidden state.
- The traditional type of feed-forward network in which all inputs feed into the first layer does not achieve this goal.

# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.

# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below



# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input  $x_t$  at each time-stamp, and a hidden state  $h_t$  that changes at each time stamp as new data points arrive.

# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input  $x_t$  at each time-stamp, and a hidden state  $h_t$  that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value  $y_t$ . For example,

# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input  $x_t$  at each time-stamp, and a hidden state  $h_t$  that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value  $y_t$ . For example,
  - In a time-series setting, the output  $y_t$  is the forecasted prediction of  $x_{t+1}$ .


# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input  $x_t$  at each time-stamp, and a hidden state  $h_t$  that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value  $y_t$ . For example,
  - In a time-series setting, the output  $y_t$  is the forecasted prediction of  $x_{t+1}$ .
  - In the text-setting, we are predicting the next word.

# Recurrent Neural Networks

- RNNs allows the input  $x_t$  to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input  $x_t$  at each time-stamp, and a hidden state  $h_t$  that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value  $y_t$ . For example,
  - In a time-series setting, the output  $y_t$  is the forecasted prediction of  $x_{t+1}$ .
  - In the text-setting, we are predicting the next word.
  - In some applications, we do not output  $y_t$  at each time stamp, but only at the end of the sequence (sentiment analysis).

# General Architecture of a RNN



# Examples of RNN

- Fully recurrent neural networks (FRNN)

# Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)



# Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)

# Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)
- Gated recurrent units (GRUs)

# The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.

# The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.

# The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.

# The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range  $(0, 1]$ , therefore, the gradient decreases exponentially with the number of layers.

# The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range  $(0, 1]$ , therefore, the gradient decreases exponentially with the number of layers.
- With activation functions whose derivatives can take on larger values, one risks encountering the related opposite exploding gradient problem.


# Table of Contents

1 Recurrent Neural Networks

2 LSTM



# Long short-term memory (LSTM)



# LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.

# LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.


# LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.

# LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.
- **LSTMs were developed to deal with the vanishing gradient problem** that can be encountered when training traditional RNNs.

## LSTM



$$i_t = \sigma(x_t U^i + h_{t-1} W^i)$$

$$f_t = \sigma(x_t U^f + h_{t-1} W^f)$$

$$o_t = \sigma(x_t U^o + h_{t-1} W^o)$$

$$\tilde{C}_t = \tanh(x_t U^g + h_{t-1} W^g)$$

$$C_t = \sigma(f_t * C_{t-1} + i_t * \tilde{C}_t)$$

$$h_t = \tanh(C_t) * o_t$$

$c_t$ : Cell state, *memoría de la célula*

$f_t$ : Forget gate, *¿qué tanto olvidamos de  $c_t$*

$o_t$ : Output gate

$i_t$ : Input gate

$i_t * \tilde{C}_t$ : *La contribución de esta entrada a la memoria*

$h_t$ : Hidden-state vector, output

# LSTM: Unfolded

