

Recurrent Neural Networks

Remembering Past Data

Dr. Mauricio Toledo-Acosta

Diplomado Ciencia de Datos con Python

Table of Contents

1 Recurrent Neural Networks

2 LSTM

Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.

Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .

Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input x_t only after the earlier inputs have already been received and converted into a hidden state.

Recurrent Neural Networks

- **Recurrent neural networks are designed for sequential data** like text sentences or time-series.
- The input is of the form x_1, \dots, x_n , where x_t is a d -dimensional point received at the time-stamp t .
- **Successive data are dependent on past data.** Therefore, it is helpful to receive a particular input x_t only after the earlier inputs have already been received and converted into a hidden state.
- The traditional type of feed-forward network in which all inputs feed into the first layer does not achieve this goal.

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .


Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .
 - In the text-setting, we are predicting the next word.

Recurrent Neural Networks

- RNNs allows the input x_t to interact directly with the hidden state created from the inputs at previous time stamps.
- The basic architecture of the recurrent neural network is illustrated below
- There is an input x_t at each time-stamp, and a hidden state h_t that changes at each time stamp as new data points arrive.
- Each time-stamp also has an output value y_t . For example,
 - In a time-series setting, the output y_t is the forecasted prediction of x_{t+1} .
 - In the text-setting, we are predicting the next word.
 - In some applications, we do not output y_t at each time stamp, but only at the end of the sequence (sentiment analysis).

General Architecture of a RNN



Examples of RNN

- Fully recurrent neural networks (FRNN)

Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)

Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)

Examples of RNN

- Fully recurrent neural networks (FRNN)
- Independently RNN (IndRNN)
- Long short-term memory (LSTM)
- Gated recurrent units (GRUs)

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.

The vanishing gradient problem

- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range $(0, 1]$, therefore, the gradient decreases exponentially with the number of layers.

The vanishing gradient problem


- The vanishing gradient problem is encountered when training artificial neural networks with gradient-based learning methods.
- In these methods, during each iteration of training each of the neural network's weights receives an update proportional to the partial derivative of the error function.
- In some cases, the gradient will be very small, preventing the weight from changing its value.
- For example, the hyperbolic tangent function have gradients in the range $(0, 1]$, therefore, the gradient decreases exponentially with the number of layers.
- With activation functions whose derivatives can take on larger values, one risks encountering the related opposite exploding gradient problem.

Table of Contents

1 Recurrent Neural Networks

2 LSTM

LSTM



$$\begin{aligned}
 i_t &= \sigma(x_t U^i + h_{t-1} W^i) \\
 f_t &= \sigma(x_t U^f + h_{t-1} W^f) \\
 o_t &= \sigma(x_t U^o + h_{t-1} W^o) \\
 \tilde{C}_t &= \tanh(x_t U^g + h_{t-1} W^g) \\
 C_t &= \sigma(f_t * C_{t-1} + i_t * \tilde{C}_t) \\
 h_t &= \tanh(C_t) * o_t
 \end{aligned}$$

c_t : Cell state, *memoría de la célula*

f_t : Forget gate, *¿qué tanto olvidamos de c_t*

o_t : Output gate

i_t : Input gate

$i_t * \tilde{C}_t$: *La contribución de esta entrada a la memoria*

h_t : Hidden-state vector, output

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.


LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.

LSTM

- Unlike standard feedforward neural networks, **LSTM has feedback connections**. LSTMs can process not only single data points, but also entire sequences of data.
- **A LSTM unit is composed of a cell, an input gate, an output gate and a forget gate**. The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell.
- LSTM networks are well-suited to classifying, and making predictions based on time series data.
- **LSTMs were developed to deal with the vanishing gradient problem** that can be encountered when training traditional RNNs.

LSTM: Unfolded



Input shape

The input to every LSTM layer must be three-dimensional.

The three dimensions of this input are:

- Samples. One sequence is one sample. A batch is comprised of one or more samples.
- Time Steps. One time step is one point of observation in the sample.
- Features. One feature is one observation at a time step.

<https://machinelearningmastery.com/reshape-input-data-long-short-term-memory-networks-keras/>

<https://stackoverflow.com/questions/60571934/which-axis-does-keras-simplernn-lstm-use-as-the-temporal-axis-by-default>